

week 10 readings

wbg231

January 2023

1 Introduction

- the second reading is like, a new york times thing to motivate why we should care about this i don't think it is needed

a model for recommendation systems

- recommendation systems have two types of entities users and items
- we get a utility matrix showing how much the user interacted with each item what ever that means in this context.
- the matrix is sparse and we interpret empty elements as meaning we have no information about the users preference on that item
- the goal is to predict the blanks in the utility matrix
- in recommendation systems we do not need to predict a users ratings of all items we just need to get the top k items for that user
- the long tail is basically in online services there are so many products to recommend so it is required to learn individual recommendations for users
- there is implicit and explicit feedback to take into account when making recommendation systems

content based recommendation systems

- the main idea of content based recommendation systems is to focus on properties of items and recommend similar items to items similar tom those we know users like
- must construct each item a profile based on characteristic's tht can be used for recommendations

- for documents can remove stop words and use the remaining words to compute tfidf scores
- we can store item features in a vector. discrete data as booleans numeric data as numeric and as long as we keep the column order consistent we can take cosine similarity
- we can also make a user profile vector which has the attributes that the user interacted with
- with both a user and item vector we can use cosine similarity to recommend items
- could also view the problem in terms of building an ml classifier model

collaborative filtering

- in collaborative filtering instead of looking at similarity of items we focus on the similarity of the users ratings for two items
- doing this requires a good distance metric
- but once we have this we can just predict any missing rating by taking the average of the n users nearest the user we are trying to predict for who have predicted this item

dimensionality reduction

- another approach to recommendation systems is to view the utility matrix as the product of some user vector and some ratings vector and try to learn this decomposition
- basically we just estimate the missing values by minimize mean square error versus what we already know